

- **Programmer: Shaun Pritchard**
- **Date: 10-10-2021**
- **Assignment: 1**
- **Prof: Michael DeGiorgio**

5:06



▼ CAP 52625 COMPUTATIONAL FOUNDATIONS OF AI

Ridge Regression using Gradient Descent - Assignment 1

Note: I decided to use symbols to make it easier to view to the implementation of code in contrast to the mathematical theorems learned in class.

Deliverables

- **Deliverable 1:** build graph of dataset $N=9$ features tuning parameter effect on inferred Ridge regression
- **Deliverable 2:** Illustrate the effect of the tuning parameter on the cross-validation error
- **Deliverable 3:** Indicate the value of λ that generated the smallest CV(5) error
- **Deliverable 4:** retrain the model of $N=400$ observations and provide the estimates of the $p=9$ best-fit model parameters.
- **Deliverable 5** Provide all your source code that you wrote from scratch to perform all analyses (aside from plotting scripts, which you do not need to turn in) in this assignment, along with instructions on how to compile and run your code.
- **Deliverable 6** Implement the assignment using statistical or machine learning libraries in a language of your choice. Compare the results with those obtained above, and provide a discussion as to why you believe your results are different if you found them to be different.

Note:

Deliverable 1-5 located here:

https://colab.research.google.com/gist/shaung1/83c9e75f7062e34897957859165f3a0d/spritchard_cap5625_programming_assignment-1_10182021.ipynb

Deliverable 6 is located at: https://colab.research.google.com/drive/1W0aaP4C2_QJo4NTQ-mrODOepyaWxnQa-?usp=sharing

▼ Import Dataset

5:06



```

1 #Math libs
2 from math import sqrt
3 from scipy import stats
4 # Data Science libs
5 import numpy as np
6 import pandas as pd
7 # Graphics libs
8 import seaborn as sns
9 import matplotlib.pyplot as plt
10 %matplotlib inline
11

```

```

1 # Mount Google Drive for data access
2 from google.colab import drive
3 drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mou



```

1 # Set up dataframe instance of dataset Credit_N400_p9.csv
2 df = pd.read_csv('/content/drive/MyDrive/Florida_Atlantic_University/Computati
3 # Set up datafeme for testing
4 # df = pd.read_csv('/content/Credit_N400_p9.csv')

```

```

1 # Check feature (row:col) shape of dataframe
2 df.shape

```

(400, 10)

```

1 # Build copy of dataset for Pre-processing
2 df1 = df.copy()
3 # Validate new dataframe
4 df1.head(3)

```

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Balance
0	14.891	3606	283	2	34	11	Male	No	Yes	333
1	106.025	6645	483	3	82	15	Female	Yes	Yes	903
2	104.593	7075	514	4	71	11	Male	No	No	580

▼ Preprocess Data

```
1 # Assign dummy variables to catigorical feature attributes
2 df1 = df1.replace({'Male': 0, 'Female':1, 'No': 0, 'Yes': 1})
```

```
1 # Validate new trianing dataframe with dummy variables
2 df1.head(3)
```

5:06



	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	B
0	14.891	3606	283	2	34	11	0	0	1	333
1	106.025	6645	483	3	82	15	1	1	1	903
2	104.593	7075	514	4	71	11	0	0	0	580

```
1 # Separate independent n X 1 feature X and convert to numpy array
2 X = np.array(df1.iloc[:, :-1], dtype='float64')
3 # Test print X feature data conversion results
4 print('Matrix shape:{X}\nValidate array:(row:col)' .format(X = X.shape), '\n')
```


```
Matrix shape:(400, 9)
Validate array:(row:col)
[[1.48910e+01 3.60600e+03 2.83000e+02 ... 0.00000e+00 0.00000e+00
 1.00000e+00]
 [1.06025e+02 6.64500e+03 4.83000e+02 ... 1.00000e+00 1.00000e+00
 1.00000e+00]
 [1.04593e+02 7.07500e+03 5.14000e+02 ... 0.00000e+00 0.00000e+00
 0.00000e+00]
 ...
 [5.78720e+01 4.17100e+03 3.21000e+02 ... 1.00000e+00 0.00000e+00
 1.00000e+00]
 [3.77280e+01 2.52500e+03 1.92000e+02 ... 0.00000e+00 0.00000e+00
 1.00000e+00]
 [1.87010e+01 5.52400e+03 4.15000e+02 ... 1.00000e+00 0.00000e+00
 0.00000e+00]]
```

```
1 # Seperate dependant n X 1 feature Y and reshape to (400 x 1) vector numpy arr
2 Y = np.array(df1.iloc[:, -1], dtype='float64').reshape([-1,1])
3 # Test print Y feature data conversion results
4 print('Dependant Feature:{Y}\n \nValidate array:(row:col)\n' .format(Y = Y.sha
5 for i in Y:
6     print(i, end = ' ')
```

```
Dependant Feature:(400, 1)
```

```
Validate array:(row:col)
```

```
[333.] [903.] [580.] [964.] [331.] [1151.] [203.] [872.] [279.] [1350.] [1407.] [0.] [2
```



▼ Center response variables and standarize features

- Convert dataframe objects to numpy arrays
- Center the response variable Y (subtracting the mean)
- Standardizing input features X to a Z score

5:06



```
1 # Center Y response variable(subtract the mean)
2 Y_p = Y- np.mean(Y, axis=0)
```

```
1 # Validate Y response vairables - mean of y_
2 y_ = np.mean(Y, axis=0)
3 print('Mean of Y:', y_)
4 print('Matrix Shape:', Y_p.shape)
```

```
Mean of Y: [520.015]
Matrix Shape: (400, 1)
```

```
1 # Split centered (row:col) of Y feature
2 Y_row, Y_col = Y_p.shape
3 # validate feature split of Y
4 print('(Y_p) Row x Col:=',Y_row, Y_col)
```

```
(Y_p) Row x Col:= 400 1
```

```
1 # Standardized X feature n x 1 matrix as X_p array and reshape
2 mean_X = np.mean(X, axis=0).reshape([1,-1])
3 # Center X
4 X_p = X - mean_X
5 # Apply standard deviation to new shape[1,-1]
6 std_X = np.std(X_p, axis=0).reshape([1,-1])
7 # Caluate centered features divided by standard deviation
8 X_p = X_p / std_X
```

```
1 # Validate X_p feature
2 print('Matrix Shape:', X_p.shape, '\n', '\n', 'Mean of X:' , '\n', mean_X, '\n'
```

```
Matrix Shape: (400, 9)
```

```
Mean of X:
[[4.5218885e+01 4.7356000e+03 3.5494000e+02 2.9575000e+00 5.5667500e+01
 1.3450000e+01 5.1750000e-01 1.0000000e-01 6.1250000e-01]]
```

```
Standard deviation of X:
[[3.52001903e+01 2.30531179e+03 1.54530616e+02 1.36955969e+00
```

```
1.72282310e+01 3.12129781e+00 4.99693656e-01 3.00000000e-01
4.87179382e-01]]
```

```
1 #Store and seperate (row:col) in variable for X_p training/test set
2 X_row , X_col = X.shape
3 print('(X_p) Row x Col:=' , X_row, X_col)
```

```
(X_p) Row x Col:= 400 9
```

5:06



▼ Assign local variables

```
1 # Local Variables
2
3 # Tuning Parmes
4  $\lambda$  = np.array([1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4])
5
6 # Learning Rate
7  $\alpha$  = 1e-5 # learning & convergence rate
8
9 # K-folds
10 k = 5
11
12 #Iterations
13 q = 1000 # iterations
14
15 #log base of lambda
16  $\lambda_{\log}$  = np.log10( $\lambda$ )
17
18 #Standardized X features
19 X_p = X_p
20
21 #Centered y features
22 Y_p = Y_p
```

Ridge Regression Batch Gradient Decent

- Implement(inferred) Batch Gradient Descent for Ridge Regression on standardizing feature X_p and centered feature Y_p numpy arrays. Where p is $n \times p$ matrix.
- α = the learning and convergence rate.
- λ = L2 regularization tuning parameters.
- q = max number of iterations (1000) as specified.
- βx = Ridge regression beta coefficients parms.

- MSE = is a storage list to contain the mean squared errors for each iteration of the Batch Gradient Descent algorithm.

Note: I decided to use symbols to make it easier to view to the implementation of code in contrast to

▼ Ridge Regression Batch Gradient Decent Algorithm

5:06



$$\beta := \beta - 2\alpha[\lambda\beta - \mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta)]$$

```

1 def RidgeRegression_BGD(X, Y, α, λ ) :
2     # Empty list to hold MSE errors
3     MSE = []
4     # Randomly initialize the parameter vector β = [β1, β2, ... , βp]
5     βx = np.random.uniform(-1,1,(X_col,1))
6     # Instantiate loop to update parameter vectors
7     for i in range(q) :
8         # Store βx coefficients in temp variable
9         βx_temp = βx
10        # Update βx parameter vector as βx := β - 2α[λβ - XT(y - Xβ)]
11        βx = βx - 2*α*( λ *βx - np.dot( X.T, Y - np.dot(X,βx) ) )
12        # Calcualte vector direction of response variables
13        ŷ = np.dot(X,βx)
14        # Instantiate temp var "MSE_temp"square-root of real Y variables
15        # minus the calculated ŷ response
16        MSE_temp = np.mean(np.square(Y - ŷ))
17        # Append updated MSE_temp caluation to MSE list
18        MSE.append(MSE_temp)
19        # Caluate the divided absolute values of βx coefficients minus βx
20        βt = np.abs((βx - βx_temp)/βx_temp)
21        # Calualte the max value of βt and store in βm
22        βm = np.max(βt)
23 # Console log to test my code:
24 #----- Feature Testing Output-----
25     if (βm < α):
26         print("Testing:\nBatch Gradient Descent(RR_BGD) breaks on: {i} ite
27         break
28     # Test for convergance error
29     if (MSE[-1] < MSE[0]): #Check MSE is lower than the initial value
30         pass
31     else :
32         print("Testing:\n Error not converging with lambda = {λ}param".format(
33         # Output updated coefficients and MSE
34         return βx, MSE

```

1 # Test tunning paremters inferred on RidgeRegression_BGD()

```

2 # Create empty list to store updated coefficients
3  $\beta\_lst = []$ 
4 # Create counter for test output
5 count = 0
6 # Iterate through RidgeRegression_BGD output:  $\beta x$ , MSE
7 for i, l in enumerate( $\lambda$ ):
8     # counter
9     count += 1
10    # print('Tuning parameters {} \n', lmbdas)
11    print('Tuning parameter converged at = #{c} $\lambda$  {} \n'.format(l, c=count,))
12    # run RidgeRegression_BGD
13     $\beta x$ , MSE = RidgeRegression_BGD(X_p, Y_p,  $\alpha$ , l)
14    # Append  $\beta x$  beta coefficients to empty list
15     $\beta\_lst.append(\beta x)$ 
16

```

5:06



Tuning parameter converged at = #1 λ 0.01

Tuning parameter converged at = #2 λ 0.1

Tuning parameter converged at = #3 λ 1.0

Tuning parameter converged at = #4 λ 10.0

Tuning parameter converged at = #5 λ 100.0

Tuning parameter converged at = #6 λ 1000.0

Testing:

Batch Gradient Descent(RR_BGD) breaks on: 418 iteration

Tuning parameter converged at = #7 λ 10000.0

Testing:

Batch Gradient Descent(RR_BGD) breaks on: 51 iteration

▼ Deliverable 1:

Build graph of dataset N=9 features tuning parameter effect on inferred Ridge regression

```

1 # Output Deliverable 1: inferred tuning parameters of ridge regression
2 sns.set_theme()
3 sns.set_style("darkgrid", {"grid.color": ".5", "image.cmap": "mako", "grid.line
4 plt.figure()
5 plt.figure(figsize=(16, 10), dpi=70)
6 plt.xlabel('λ Tuning Params')
7 plt.ylabel('p=9 features')

```

```


8 plt.title('Deliverable 1: Inferred Ridge Regression Coefficients')
9 for i in range(X_col) :
10      $\beta_j = [\beta_x[i,0] \text{ for } \beta_x \text{ in } \beta\_lst ]$ 
11     legend = 'Beta_λ_{}'.format(i)
12     sns.lineplot(x=λ_log, y=βj , label=legend, )
13 # Output Deliverable1.jpg to file
14 plt.savefig('SPritchard_CAP5625_Assignment1_Deliverable1.jpg')
15 plt.show()
16
17

```

5:06



<Figure size 432x288 with 0 Axes>



▼ (5)K-Fold Grid Search Cross Validation Algorithm

- (5)K-fold Grid Cross Validation calualted Batch Gradient Decent Ridge Regression with hyperparameter tuning

- Use grid search CV to train and test 5 k folds
- Calculate test and training errors
- Test MSE on tuning params

5:06



```

1 # Implement start and end of test k-fold data split
2 X_row = X.shape[0]
3 # Divide absolute row features by k = 5
4 X_row_test = X_row // k
5 # Store data in k-fold array (Kfold/Kfold_ )
6 Kfold = [ X_row_test * ind for ind in range(k)] # initial k-folds
7 Kfold_ = [ ind + X_row_test for ind in Kfold ] # End k-folds

1 # Grid Cross Validation for Batch Gradient Decent Ridge Regression with hyperp
2 # Instantiate empty list to hold Cross vladations errors
3 CV = []
4 # Add a counter to iterate tuning params and errors
5 for i, l in enumerate( $\lambda$ ) :
6     # print('(5)K-fold CV tuning parameter error = {}'.format(l))
7     MSE = []
8     # Loop through K trianing and test vectors
9     for i in range(k):
10         #Hold-out 5 k-folds arrays (80 x 9)
11         CV_fold = Kfold[i]
12         CV_fold_ = Kfold_[i]
13         # Seperate training feature variables
14         X_train = np.row_stack(( X[0:CV_fold,:] , X[CV_fold_:, :] ))
15         Y_train = np.row_stack(( Y[0:CV_fold,:] , Y[CV_fold_:, :] ))
16         # Seperate testing feature variables
17         X_test = X[CV_fold:CV_fold_, :]
18         Y_test = Y[CV_fold:CV_fold_, :]
19         # Standardize X test set
20         X_test_ = (X_test - np.mean(X_test, axis=0))/np.std(X_test, axis=0)
21         # Center Y test set
22         Y_test_ = Y_test - np.mean(Y_test, axis=0)
23         # Implement ridge regressionand MSE on test data
24          $\beta_x$ , _ = RidgeRegression_BGD(X_p, Y_p,  $\alpha$ , l)
25         # Product transofrmation of test data on trining set
26          $\hat{y}$  = np.dot(X_test_,  $\beta_x$ )
27         # Claulate average squareroot of Y(test)-  $\hat{y}$  variables
28         err = np.mean(np.square(Y_test_ -  $\hat{y}$ ))
29         # Append calualtion to MSE list
30         MSE.append(err)
31         #Caluate average of updated MSE
32         CV_err = np.mean(MSE)

```

```

33     # Append averaged MSE variables to CV list
34     CV.append(CV_err)

```

```

Testing:
Batch Gradient Descent(RR_BGD) breaks on: 420 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 423 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 428 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 422 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 427 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 55 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 54 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 51 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 55 iteration
Testing:
Batch Gradient Descent(RR_BGD) breaks on: 49 iteration

```

5:06



```

1 # console log  to test CV code
2 #----- Feature Testing -----
3 print('Initial (5)K-folds', Kfold, '\n')
4 print('Ending (5)K-folds', Kfold_, '\n')
5 print(X_test.shape, Y_test.shape, X_train.shape, Y_train.shape, '\n')
6 print(X_test[0,0], Y_test[0,0], X_train[0,0], Y_train[0,0], '\n')
7 print("Mean Square Error",err, '\n')
8 print("MSE value",MSE, '\n')
9 print("CV_error", CV_err, '\n')
10

```

Initial (5)K-folds [0, 80, 160, 240, 320]

Ending (5)K-folds [80, 160, 240, 320, 400]

(80, 9) (80, 1) (320, 9) (320, 1)

16.279 5.0 14.890999999999998 333.0

Mean Square Error 167309.80185008282

MSE value [167455.13745409623, 220062.24519126484, 192200.85842315012, 159943.406401254

CV_error 181394.28986396975



▼ Deliverable 2

- Illustrate the effect of the tuning parameter on the cross validation error by generating a plot.
- Grid Cross Validation tuning errors for each tuning parameter value, perform five-fold cross validation and choose the value of λ that gives the smallest value.


5:06



$$CV_{(5)} = \frac{1}{5} \sum_{i=1}^5 MSE_i$$

```
1 # Illustrate the effect of the tuning parameter on the cross validation error
2 sns.set_theme()
3 sns.set_style("darkgrid", {"grid.color": ".5", "image.cmap": "mako", "grid.line
4 plt.figure()
5 plt.figure(figsize=(16, 10), dpi=70)
6 plt.title('Deviverable 2: Effect of tuning parameter on the Cross Validation E
7 plt.xlabel('λ Tuning Params')
8 plt.ylabel('CV Error')
9 sns.lineplot(x=λ_log, y=CV , color='purple', markersize=12)
10 plt.savefig('SPritchard_CAP5625_Assignment1_Deliverable2.jpg')
11 plt.show()
```

<Figure size 432x288 with 0 Axes>



▼ Deliverable 3

- Indicate the value of λ that generated the smallest CV(5) error

```

1 #Find minimum MSE error
2 err = min(MSE)
3 # index MSE error
4 i = MSE.index(err)
5 # Iterate to find MSe from  $\lambda$  list
6 l =  $\lambda$ [i]
7 # Output final results of lowest  $\lambda$  tuning param
8 print("Best CV error of  $\lambda$  = {e}\nBest tuning param of  $\lambda$  = {l}".format(e=err,
```

```

Best CV error of  $\lambda$  = 159943.40640125488
Best tuning param of  $\lambda$  = 10.0
```

▼ Deliverable 4

- Given the optimal λ , retrain your model on the entire dataset of $N = 400$ observations and provide the estimates of the $p = 9$ best-fit model parameters.

```

1 # Retrain model based on  $\lambda = 10.0$ 
2  $\beta x$ , _dh = RidgeRegression_BGD(X_p, Y_p,  $\alpha$ , l)
3 # Output best fit model params of  $\beta x$  based on on  $\lambda = 10.0$  tuning param
4 print('Best fit model parameters', '\n',  $\beta x$ )
```

```

Best fit model parameters
[[-206.6352321 ]
 [ 277.44127207]
 [ 273.40973826]
 [ 22.38621464]
 [-17.61875517]
 [-1.74246626]
 [-3.0415899 ]
 [ 121.65556358]
 [-5.61633766]]
```

5:06

▶

🔖

✓ 0s completed at 1:43 PM

● ✕