

- **Programmer:**
 - Shaun Pritchard
 - Ismael A Lopez

▼ Assignment 3

Brief overview of assignment

perform a penalized (regularized) logistic (multinomial) regression fit using ridge regression, with the model parameters obtained by batch gradient descent. Your predictions will be based on K=5 continental ancestries (African, European, East Asian, Oceanian, or Native American). Ridge regression will permit you to provide parameter shrinkage (tuning parameter $\lambda=0$) to mitigate overfitting. The tuning parameter λ will be chosen using five-fold cross validation, and the best-fit model parameters will be inferred on the training dataset conditional on an optimal tuning parameter. This trained model will be used to make predictions on new test data points

Table of Contents

- [Import Packages](#)
 - [Import packages for manipulating data](#)
 - [Import packages for splitting data](#)
 - [Import packages for modeling data](#)
 - [Import packages for Scaling and Centering data](#)
 - [Import packages for Measuring Model Performance](#)
- [Data Processing](#)
 - [Import Data](#)
 - [Let's change the categorical values](#)
 - [Create Predictor and Target numpy array](#)
 - [Create a Normalize copy of variables](#)
 - [Split Data](#)
- [Regression Model](#)
 - [Define our learning rates:](#)
 - [Create the Regression Objects](#)
 - [LogisticRegression Library](#)

Deliverables

- [Deliverable 6.1](#)
- [Deliverable 6.2](#)
- [Deliverable 6.3](#)
- [Deliverable 6.4](#)
- [Deliverable 6. Reason for difference](#)

▼ Import Packages

▼ Import packages for manipulating data

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import matplotlib as mpl
5 import matplotlib.mlab as mlab
6 import math
7 import csv
8 import random
9 %matplotlib inline
10 from sklearn.preprocessing import LabelEncoder
11 import seaborn as sns
12 import math
13
```

▼ Import packages for splitting data

```
1 from sklearn.model_selection import train_test_split, cross_val_score, KFold,
2 from sklearn.model_selection import GridSearchCV
3
```

▼ Import packages for modeling data

```
1 # Import models:
2 from sklearn.preprocessing import PolynomialFeatures
3 from sklearn.linear_model import LinearRegression as linearR_Model, Ridge as
4 from sklearn.linear_model import RidgeCV
5
```

```
6 from sklearn.linear_model import ElasticNet
7 from sklearn.linear_model import ElasticNetCV
8
9 from sklearn.linear_model import LogisticRegression
10
11
12 from sklearn.exceptions import ConvergenceWarning
13 #from sklearn.utils._testing import ignore_warnings
14 import warnings
15 warnings.filterwarnings('ignore', category=ConvergenceWarning) # To filter out
16 warnings.filterwarnings('ignore', category=UserWarning)
17 from itertools import product
18
```

▼ Import packages for Scaling and Centering data

```
1 from sklearn.preprocessing import StandardScaler
```

▼ Import packages for Measuring Model Perormance

```
1 from sklearn.metrics import mean_squared_error
2 from sklearn.metrics import r2_score
3 from sklearn.metrics import make_scorer
```

▼ Data Processing

▼ Import Data

Traing Dataset

```
1 Train_dataset = pd.read_csv ('TrainingData_N183_p10.csv')
2 Train_dataset.head(3)
```

PC1 PC2 PC3 PC4 PC5 PC6 PC7 PC8

```
1 # What are the datatypes of each observation:
2 print(Train_dataset.dtypes)
3 # Shape of my data
4 print('The size of our data are: ',Train_dataset.shape)
```

```
PC1      float64
PC2      float64
PC3      float64
PC4      float64
PC5      float64
PC6      float64
PC7      float64
PC8      float64
PC9      float64
PC10     float64
Ancestry  object
dtype: object
The size of our data are: (183, 11)
```

```
1 print('Training Dataset Missing Values: \n',Train_dataset.isnull().sum())
2
```

```
Training Dataset Missing Values:
PC1      0
PC2      0
PC3      0
PC4      0
PC5      0
PC6      0
PC7      0
PC8      0
PC9      0
PC10     0
Ancestry  0
dtype: int64
```

Test Dataset

```
1 Test_dataset = pd.read_csv ('TestData_N111_p10.csv')
2 Test_dataset.head(3)
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	
0	2.517683	5.464283	9.067873	-4.965928	-0.741937	0.039785	0.573279	-0.216918	2.45
1	6.077012	1.032867	-5.795883	-3.490064	-0.600204	-0.120803	1.243767	1.821390	-1.17
2	1.016945	-2.913299	0.907702	1.233580	-1.983452	1.605964	2.674998	-0.732921	-2.15

```

1 # What are the datatypes of each observation:
2 print(Test_dataset.dtypes)
3 # Shape of my data
4 print('The size of our data are: ',Test_dataset.shape)

```

```

PC1          float64
PC2          float64
PC3          float64
PC4          float64
PC5          float64
PC6          float64
PC7          float64
PC8          float64
PC9          float64
PC10         float64
Ancestry     object
dtype: object
The size of our data are: (111, 11)

```

```

1 # Are there any null or missing values
2 print('Test Dataset Missing Values: \n',Test_dataset.isnull().sum())

```

```

Test Dataset Missing Values:
PC1          0
PC2          0
PC3          0
PC4          0
PC5          0
PC6          0
PC7          0
PC8          0
PC9          0
PC10         0
Ancestry     0
dtype: int64

```

▼ Lets change the categorical values

```

1 # recode the categories
2 Training_Class = Train_dataset['Ancestry'].unique().tolist()
3 Test_Class = Test_dataset['Ancestry'].unique().tolist()
4 num_features = len(Training_Class)
5
6
7 print("Unique Values for Train Ancestry: ", Training_Class)
8 print("Unique Values for Test Ancestry: ", Test_Class)
9

```

Unique Values for Train Ancestry: ['African', 'European', 'EastAsian', 'Oceanian', 'Na
 Unique Values for Test Ancestry: ['Unknown', 'Mexican', 'AfricanAmerican']

```
1 from sklearn.preprocessing import LabelEncoder
2 le = LabelEncoder()
3 Train_dataset['Ancestry_Encoded'] = le.fit_transform(Train_dataset.iloc[:, -1:])
4 Test_dataset['Ancestry_Encoded'] = le.fit_transform(Test_dataset.iloc[:, -1:])
```

```
1 Train_dataset.head(3)
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
0	-10.901171	0.798743	-1.143301	-1.070960	11.856396	-2.265965	4.536405	1.519959
1	-9.990054	1.416821	-0.729626	-0.443621	10.418594	0.443514	2.640659	-4.637746
2	-9.345388	2.913054	-0.921421	0.029173	10.672615	-2.052552	5.140476	-1.451096

▼ Create Predictor and Target numpy array

```
1 # Target:
2 Y_Train= Train_dataset['Ancestry_Encoded'].to_numpy()
3 Y_Test= Test_dataset['Ancestry_Encoded'].to_numpy()
4 Y_Train.shape
```

```
(183,)
```

```
1 # Convert the Pandas dataframe to numpy ndarray for computational improvement
2 X_Train = Train_dataset.iloc[:, :-2].to_numpy()
3 X_Test = Test_dataset.iloc[:, :-2].to_numpy()
4
5 print(type(X_Train), X_Train[:1], "Shape = ", X_Train.shape)
```

```
<class 'numpy.ndarray'> [[-10.90117144  0.79874334 -1.14330096 -1.07096001 11.85639
-2.2659654  4.5364047  1.51995913 -2.21429419 -0.67127393]] Shape = (183, 10)
```

▼ Create a Normalize copy of variables

```
1 # Create Standarizing ObjectPackages:
2 standardization = StandardScaler()
3
4 # Strandardize
```

```

5 n_observations = len(Train_dataset)
6 variables = Train_dataset.columns
7
8
9 # Standardize the Predictors (X)
10 X_Train = standardization.fit_transform(X_Train)
11
12 # Add a constant to the predictor matrix
13 X_Train = np.column_stack((np.ones(n_observations),X_Train))
14
15
16 # Save the original M and Std of the original data. Used for unstandardize
17 original_means = standardization.mean_
18
19 # we changed standardization.std_ to standardization.var_**.5
20 originanal_stds = standardization.var_**.5
21
22
23 print("observations :", n_observations)
24 print("variables :", variables[:2])
25 print('original_means :', original_means)
26 print('originanal_stds :', originanal_stds)
27
28
29

```

```

observations : 183
variables : Index(['PC1', 'PC2'], dtype='object')
original_means : [ 1.40487976e+00  2.02293488e+00  1.91271130e-03  1.02811502e-01
 2.43929372e-01  2.93901516e-01  4.37620184e-02 -1.85769325e-01
 1.03879526e-01 -4.17198356e-02]
originanal_stds : [4.8993287  3.47654999 3.90903976 3.149965  2.14032401 1.77048761
 1.58593444 1.50391174 1.58141009 0.97706561]

```

Split Data:

let's first split it into train and test part

```
X_train, X_out_sample, y_train, y_out_sample = train_test_split(Xst, y_Centered, test_size=0.40,
random_state=101) # Training and testing split
```

```
X_validation, X_test, y_validation, y_test = train_test_split(X_out_sample, y_out_sample,
test_size=0.50, random_state=101) # Validation and test split
```

Print Data size

```
print ("Train dataset sample size: {}".format(len(X_train))) print ("Validation dataset sample size:
0" format(len(X_validation))) print ("Test dataset sample size: 0" format(len(X_test)))
```

▼ Regression Model

▼ Define our learning rates

```
1 # Define my tuning parameter values  $\lambda$ :
2
3 learning_rates_λ = [1.e-04, 1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+02, 1.
4 print(learning_rates_λ)
5
6 # learning rate
7 α = 1e-4
8
9 # K-folds
10 k = 5
11
12
13 # Itterations
14 n_iters = 10000
15
```

[0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]

▼ Create the Regression Objects

LogisticRegression Library

```
1 # LogisticRegression
2 from sklearn.linear_model import LogisticRegression
3 Library_LogisticRegression = LogisticRegression(max_iter = 10000, multi_clas:
4
```

▼ Deliverable 7.1

Deliverable 1: Illustrate the effect of the tuning parameter on the inferred ridge regression coefficients by generating five plots (one for each of the $K=5$ ancestry classes) of 10 lines (one for each of the $p=10$ features), with the y -axis as β_{jk} , $j=1,2,\dots,10$ for the graph of class k , and x -axis the corresponding log-scaled tuning parameter value $\log_{10}(\lambda)$ that 7 generated the particular β_{jk} . Label both axes in all five plots. Without the log scaling of the tuning parameter, the plot will look distorted.

LogisticRegression with Library

```

1 Lβ_per_λ=[] # set empty list
2
3 # Evaluate tuning parameters with LogisticRegression penalty
4 for tuning_param in learning_rates_λ:
5     Library_LogisticRegression = LogisticRegression(max_iter = 10000, mult
6     Library_LogisticRegression.fit(X_Train, Y_Train)
7     c = np.array(Library_LogisticRegression.coef_)
8     # c = np.append(tuning_param,c)
9     Lβ_per_λ.append(Library_LogisticRegression.coef_)
10 #     print(c)
11

```

```

1 Lβ_per_λ[0]

array([[ -5.99624844e-03,  2.96117039e-04, -8.73802214e-04,
        -2.04099800e-04,  1.39486513e-03, -5.71819264e-04,
        -5.20972427e-04,  1.07976342e-04,  1.04013395e-03,
        -6.51753379e-04],
       [ 3.17782901e-03,  3.20405131e-03, -2.75595421e-04,
        7.25605495e-03, -5.20324070e-04, -9.23717795e-04,
        -9.90076526e-05,  7.45094436e-04,  4.01783544e-04,
        -1.52613597e-03],
       [ 1.95202263e-04, -6.24913113e-03,  9.13687544e-04,
        -4.18520330e-04, -5.83513899e-04,  3.02492530e-03,
        7.84272369e-04, -2.65956521e-03, -1.24063618e-03,
        -4.04365901e-04],
       [ 2.28441492e-03, -5.37970381e-04, -4.83132244e-03,
        -3.44509733e-03, -1.70996181e-04, -1.31009775e-03,
        -9.47052680e-05,  1.90245786e-03,  5.30506344e-04,
        3.67778162e-03],
       [ 3.38802248e-04,  3.28693316e-03,  5.06703253e-03,
        -3.18833750e-03, -1.20030982e-04, -2.19290489e-04,
        -6.95870212e-05, -9.59634295e-05, -7.31787661e-04,
        -1.09552637e-03]])

```

```

1 # Loop throught the betas, by class generated by each lamda
2 temp_df = []
3 for l in range(np.array(Lβ_per_λ).shape[0]):

```

```

4     for c in range(np.array(L $\beta$ _per_ $\lambda$ ).shape[1]):
5         temp_df.append(np.append(L $\beta$ _per_ $\lambda$ [1][c],(learning_rates_ $\lambda$ [1],c)))

1 TunnedL $\beta$ _df=pd.DataFrame(np.array(temp_df))
2 TunnedL $\beta$ _df.columns=['PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6', 'PC7', 'PC8', '
3 #TunnedL $\beta$ _df['Class_Name'] = TunnedL $\beta$ _df['Class_Name'].apply(lambda x: Trainin
4 TunnedL $\beta$ _df.head(10)

```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	
0	-0.005996	0.000296	-0.000874	-0.000204	0.001395	-0.000572	-0.000521	0.000108	0.0
1	0.003178	0.003204	-0.000276	0.007256	-0.000520	-0.000924	-0.000099	0.000745	0.0
2	0.000195	-0.006249	0.000914	-0.000419	-0.000584	0.003025	0.000784	-0.002660	-0.0
3	0.002284	-0.000538	-0.004831	-0.003445	-0.000171	-0.001310	-0.000095	0.001902	0.0
4	0.000339	0.003287	0.005067	-0.003188	-0.000120	-0.000219	-0.000070	-0.000096	-0.0
5	-0.058225	0.002975	-0.008652	-0.001678	0.013273	-0.005503	-0.005145	0.001171	0.0
6	0.030222	0.029984	-0.002905	0.068954	-0.004819	-0.008493	-0.000758	0.006758	0.0
7	0.002223	-0.059828	0.008914	-0.003635	-0.005770	0.028374	0.007473	-0.024838	-0.0
8	0.022178	-0.004906	-0.046279	-0.032885	-0.001479	-0.012364	-0.000878	0.017768	0.0
9	0.003602	0.031775	0.048922	-0.030756	-0.001205	-0.002014	-0.000692	-0.000860	-0.0

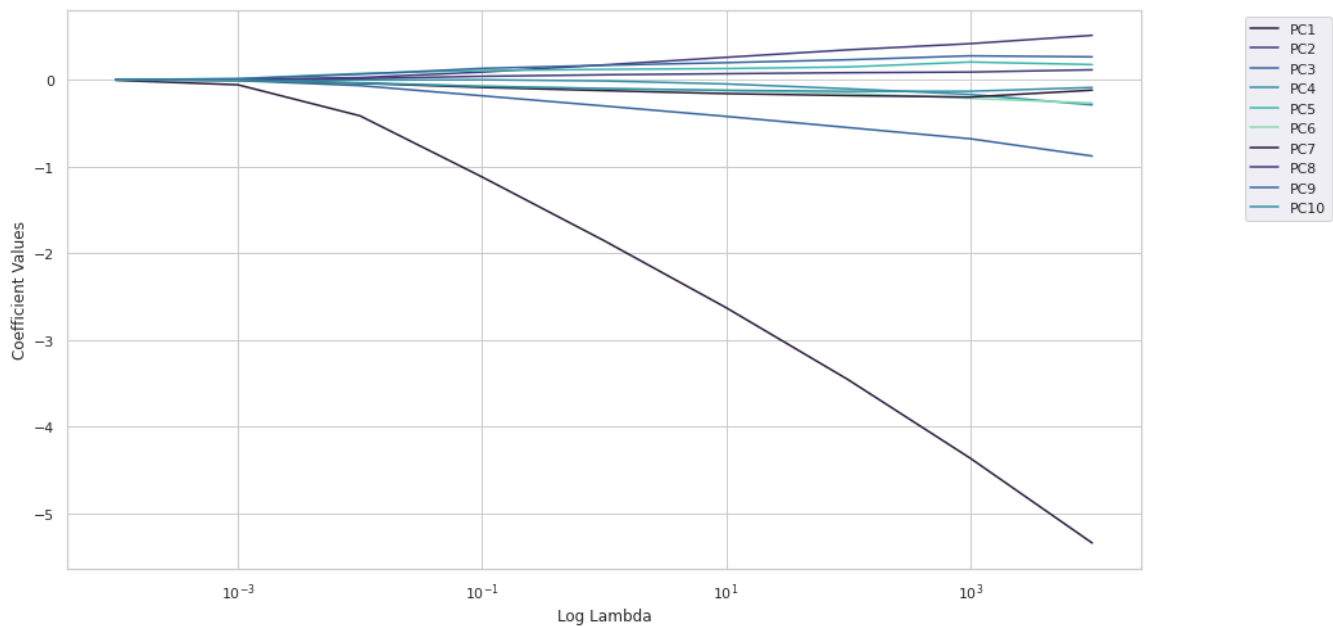
```
1 Training_Class
```

```
['African', 'European', 'EastAsian', 'Oceanian', 'NativeAmerican']
```

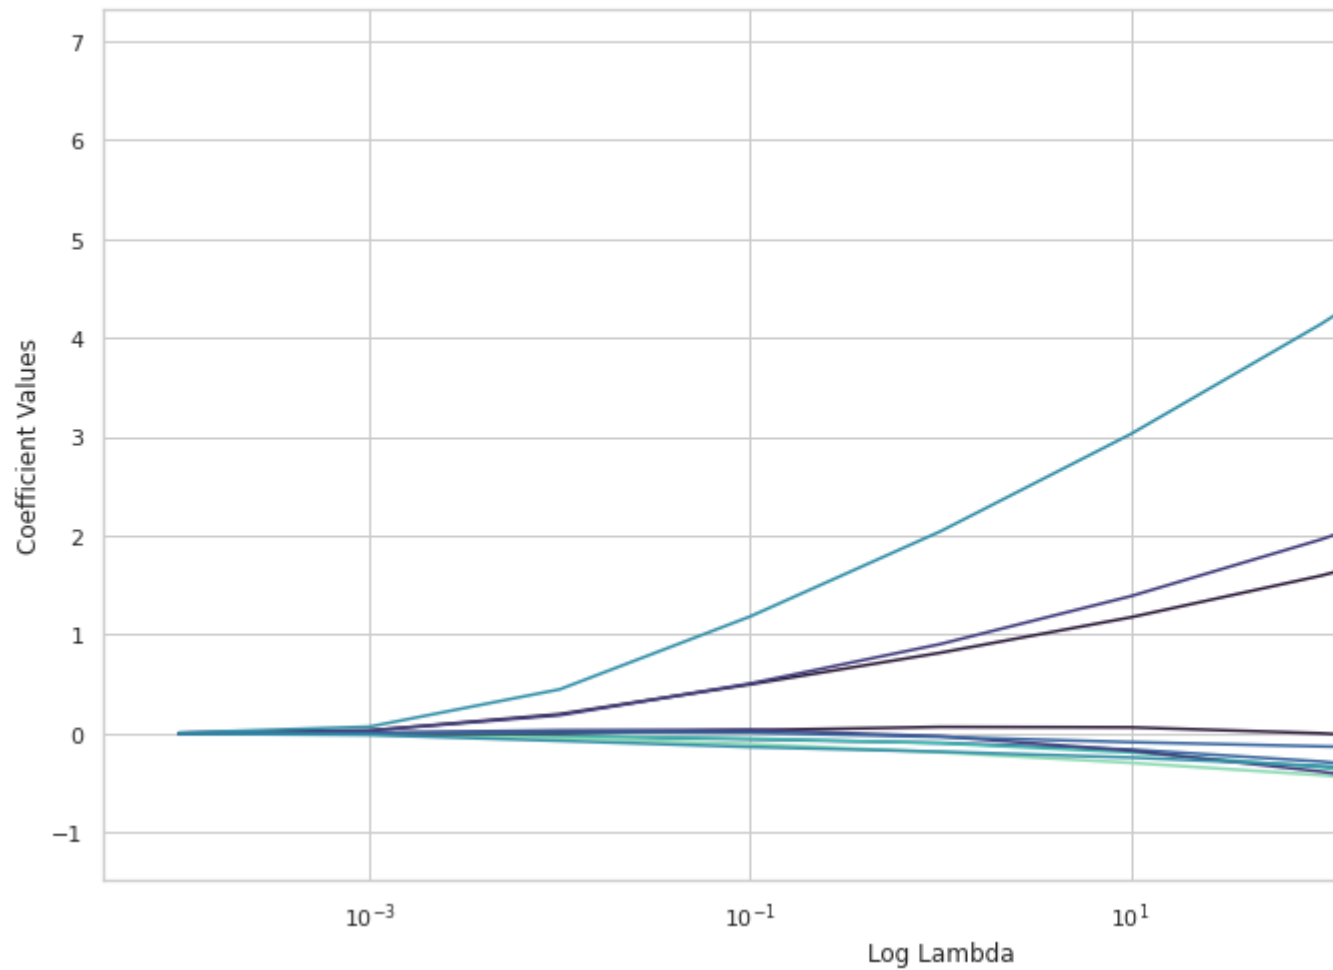
```
1 TunnedL $\beta$ _df[TunnedL $\beta$ _df.Class.eq(0)]
```

```
1 # Plot tuning parameter on the inferred ridge regression coefficients
2 sns.set(rc = {'figure.figsize':(15,8)})
3 for i, c in enumerate(Training_Class):
4     sns.set_theme(style="whitegrid")
5     sns.set_palette("mako")
6     for j in range(1, 1 + X_Train.shape[1]):
7         sns.lineplot( x = TunnedLβ_df[TunnedLβ_df.Class.eq(i)][ 'Lambda'], y =
8             sns.set()
9             plt.xscale('log')
10            plt.legend(bbox_to_anchor=(1.09, 1), loc='upper left')
11            plt.xlabel('Log Lambda')
12            plt.ylabel('Coefficient Values')
13            plt.suptitle('Inferred Ridge Regression Coefficient Tuning Parameters of'
14            plt.show()
15
```

Inferred Ridge Regression Coefficient Tuning Parameters of African Class


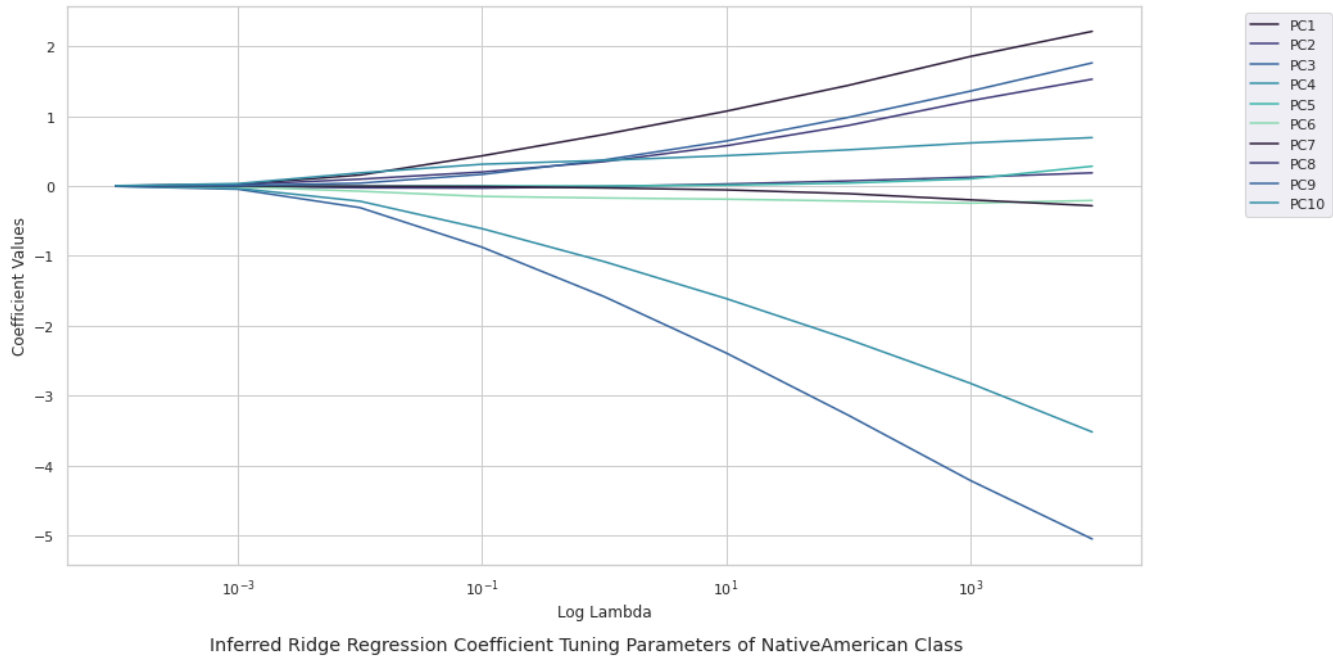
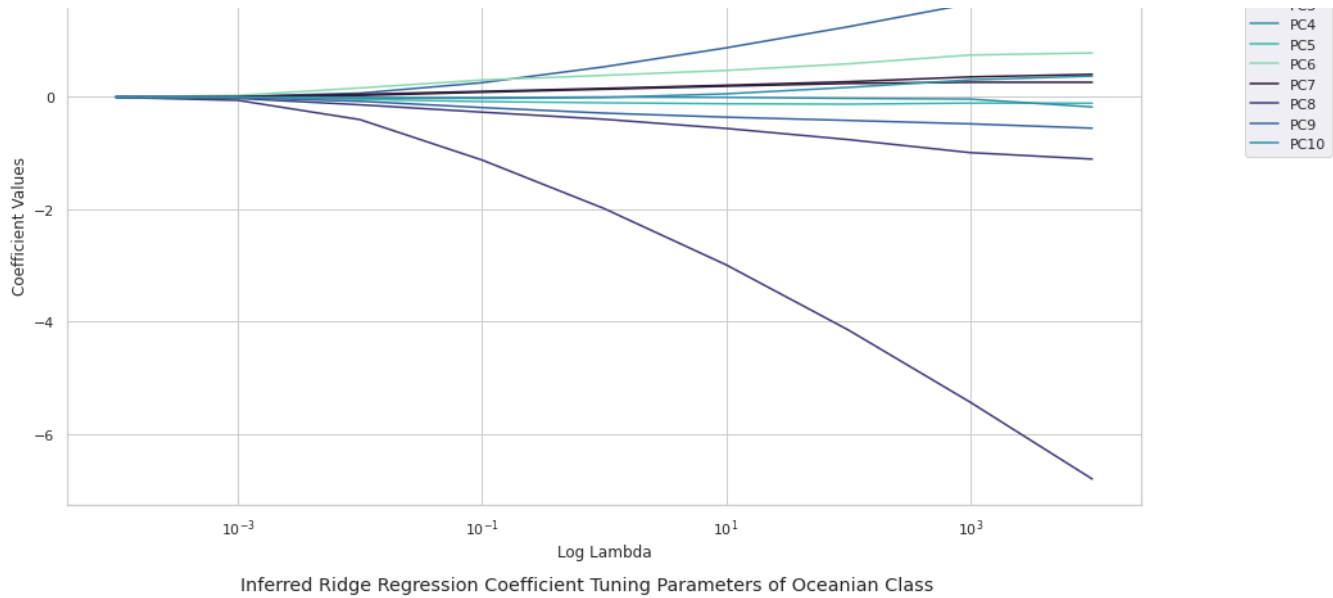


Inferred Ridge Regression Coefficient Tuning Parameters of European Class



Inferred Ridge Regression Coefficient Tuning Parameters of EastAsian Class





▼ Deliverable 7.2

Illustrate the effect of the tuning parameter on the cross validation error by generating a plot with the y -axis as CV(5) error, and the x -axis the corresponding log-scaled tuning parameter value $\log_{10}(\lambda)$ that generated the particular CV(5) error. Label both axes in the plot. Without the log scaling of the tuning parameter λ , the plots will look distorted.

CV Elastic Net with Library

```

1  from sklearn.model_selection import GridSearchCV
2  from sklearn.linear_model import LogisticRegression
3
4  #Define the model
5  Library_LogisticRegression = LogisticRegression(max_iter=10000, multi_class='ovr')
6
7
8  # Create the Kfold:
9  cv_iterator = KFold(n_splits = 5, shuffle=True, random_state=101)
10
11 cv_score = cross_val_score(Library_LogisticRegression, X_Train, Y_Train, cv=cv_iterator)
12 print (cv_score)
13 print ('Cv score: mean %0.3f std %0.3f' % (np.mean(cv_score), np.std(cv_score)))
14
15
16 [-0. -0. -0. -0. -0.]
17 Cv score: mean 0.000 std 0.000
18
19
20 1 # define grid
21 2 Parm_grid = dict()
22 3 Parm_grid['C'] = learning_rates_lambda
23 4 Parm_grid

```

```
{'C': [0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]}
```

```
1 # Lets define search
```

```
2 GsearchCV = GridSearchCV(estimator = Library_LogisticRegression, param_grid =
```

```
3 GsearchCV.fit(X_Train, Y_Train)
```

```
4
```

```
5
```

```
GridSearchCV(cv=KFold(n_splits=5, random_state=101, shuffle=True),
              estimator=LogisticRegression(max_iter=10000,
                                           multi_class='multinomial'),
              n_jobs=1,
              param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0,
                                1000.0, 10000.0]},
              scoring='neg_mean_squared_error')
```

```
1 GCV_df = pd.concat([pd.DataFrame(GsearchCV.cv_results_["params"]),pd.DataFrame
```

```
2 #GCV_df.index=GCV_df['alpha']
```

```
3 GCV_df.rename(columns={"C": "learning_rates_λ"}, inplace=True)
```

```
4
```

```
5 GCV_df
```

	learning_rates_λ	mean_test_score
0	0.0001	-2.404354
1	0.0010	-2.404354
2	0.0100	-0.322222
3	0.1000	0.000000
4	1.0000	0.000000
5	10.0000	0.000000
6	100.0000	0.000000
7	1000.0000	0.000000
8	10000.0000	0.000000

```
1
```

```
2 sns.set_theme(style="whitegrid")
```

```
3 sns.set_palette("mako")
```

```
4
```

```
5
```

```
6 plt.plot(GCV_df["learning_rates_λ"] , GCV_df["mean_test_score"])
```

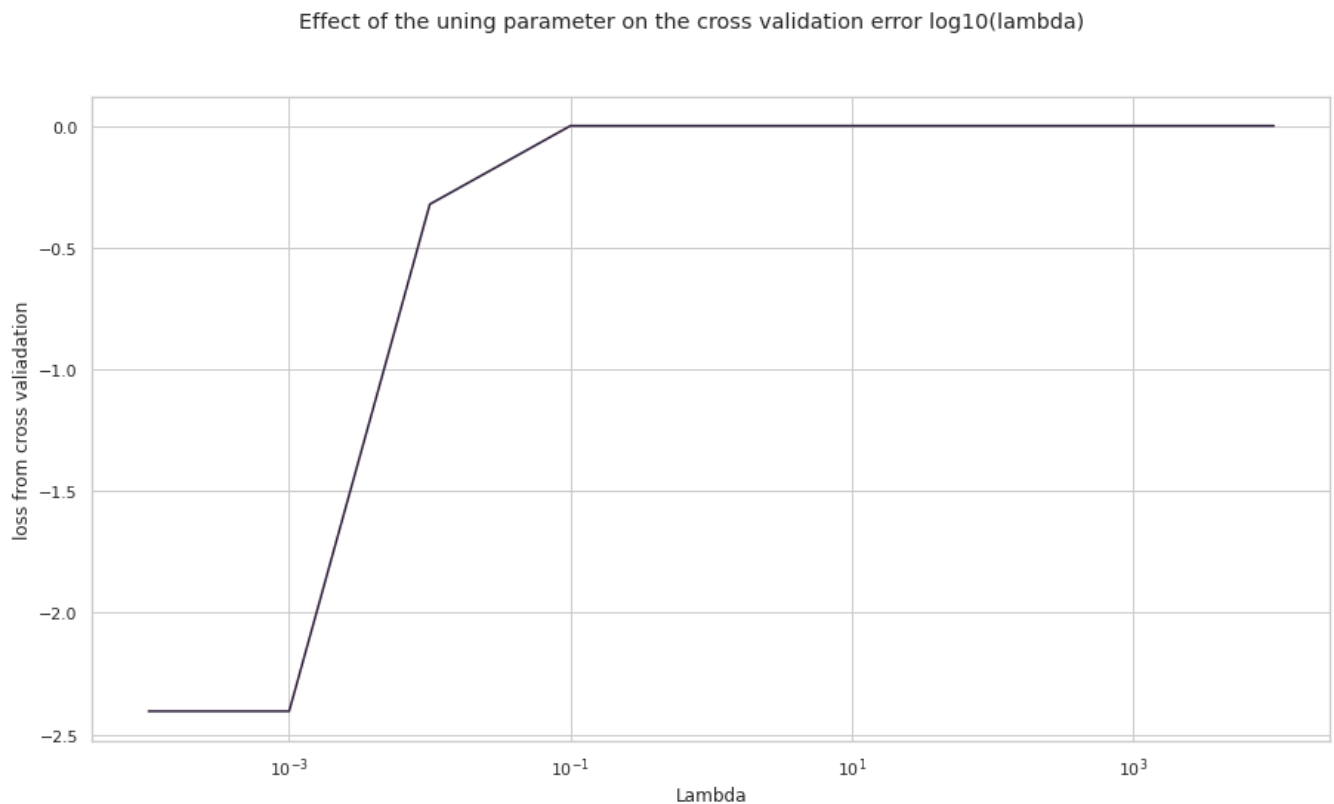
```
7
```

```
8
```

```

9 sns.set_palette("mako")
10 sns.set()
11
12
13 plt.suptitle('Effect of the uning parameter on the cross validation error log1
14 plt.xscale('log')
15 plt.xlabel('Lambda')
16 plt.ylabel('loss from cross valiadation')
17 plt.show()

```



▼ Deliverable 7.3

Indicate the value of λ that generated the smallest CV(5) error

Smallest CV with Library

```
1 print ('Best: ', GsearchCV.best_params_)
```



```
2 print ('Best CV mean squared error: %0.3f' % np.abs(GsearchCV.best_score_))
```

```
Best: {'C': 0.1}
```

```
Best CV mean squared error: 0.000
```

```
1 GCV_df.sort_values(by=['mean_test_score'], ascending=False)[:1]
```

```
2
```

	learning_rates_λ	mean_test_score
3	0.1	0.0

```
1 # Alternative: sklearn.linear_model.ElasticNetCV
```

```
2 from sklearn.linear_model import LogisticRegressionCV
```

```
3
```

```
4 auto_LR = LogisticRegressionCV(Cs = learning_rates_λ, cv=5, max_iter = 10000,
```

```
5 auto_LR.fit(X_Train, Y_Train)
```

```
6 #print ('Best alpha: %0.5f' % auto_LR.alpha_)
```

```
7 print ('Best λ: ', auto_LR.C_)
```

```
Best λ: [0.1 0.1 0.1 0.1 0.1]
```

▼ Deliverable 7.4

Given the optimal λ , retrain your model on the entire dataset of $N=183$ observations to obtain an estimate of the $(p+1) \times K$ model parameter matrix as \mathbf{B} and make predictions of the probability for each of the $K=5$ classes for the 111 test individuals located in TestData_N111_p10.csv. That is, for class k , compute $p_k(X; \mathbf{B}) = \exp(\beta_0 k + \sum X_j \beta_{jk}) / \sum \exp(\beta_0 \ell + \sum X_j \beta_{j\ell})$

- for each of the 111 test samples X , and also predict the most probable ancestry label as
 - $Y(X) = \arg \max_{k \in \{1, 2, \dots, K\}} p_k(X; \mathbf{B})$
- Report all six values (probability for each of the $K=5$ classes and the most probable ancestry label) for all 111 test individuals.

Tunned with best λ with Library

```
1 Library_LogisticRegression_best= LogisticRegression(max_iter = 10000, multi_
```

```
2 Library_LogisticRegression_best.fit( X_Train, Y_Train )
```

```
3
```

```
4 y_predM_best = Library_LogisticRegression_best.predict(X_Test)
```

```
5 print ("Betas= ", np.mean(Library_LogisticRegression_best.coef_, 0))
```

```
6
```

```
7
```

```
Betas= [ 1.97064587e-16  1.11022302e-16 -4.44089210e-17  0.00000000e+00
 -5.41233725e-17 -6.03683770e-17 -1.94289029e-17  2.25514052e-17
 -4.99600361e-17  4.99600361e-17]
```

```
1 yhat = Library_LogisticRegression_best.predict_proba(X_Test)
2 # summarize the predicted probabilities
3 print('Predicted Probabilities: %s' % yhat[0])
```

```
Predicted Probabilities: [2.65398539e-08 4.07623241e-07 4.19578741e-08 3.58579025e-08
 9.99999488e-01]
```

```
1  $\hat{y}_{test}$  = Library_LogisticRegression_best.predict_proba(X_Test)
2  $\hat{y}_{test}[:3]$ 
```

```
array([[2.65398539e-08, 4.07623241e-07, 4.19578741e-08, 3.58579025e-08,
        9.99999488e-01],
       [3.42825046e-08, 3.12778397e-05, 2.56848531e-06, 9.99964053e-01,
        2.06669173e-06],
       [3.56852016e-04, 2.09402026e-02, 9.76706428e-01, 7.64024238e-04,
        1.23249309e-03]])
```

```
1 Y_class = Library_LogisticRegression_best.predict(X_Test)
2 Y_class
```

```
array([4, 3, 2, 0, 1, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 2, 2, 1,
        3, 3, 3, 3, 3, 1, 3, 2, 2, 2, 3, 3, 2, 3, 3, 2, 3, 3, 3, 3, 3,
        3, 3, 3, 3, 3, 3, 2, 2, 2, 3, 3, 3, 3, 4, 3, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0])
```

```
1 # Re-label feature headers and add new class prediction index column
2 new_colNames = ['{}_Probability'.format(c_name) for c_name in Training_Class]
3 new_colNames
```

```
['African_Probability',
 'European_Probability',
 'EastAsian_Probability',
 'Oceanian_Probability',
 'NativeAmerican_Probability',
 'ClassPredInd']
```

```
1 # Implement index array of probabilities
2 i_prob = np.concatenate(( $\hat{y}_{test}$ , Y_class[:, None]), 1)
```

```
1 # Create New dataframe for probability indices
2 df2 = pd.DataFrame(i_prob, columns = new_colNames)
```

3 df2

	African_Probability	European_Probability	EastAsian_Probability	Oceanian_Probability
0	2.653985e-08	4.076232e-07	4.195787e-08	3.58579
1	3.428250e-08	3.127784e-05	2.568485e-06	9.99964
2	3.568520e-04	2.094020e-02	9.767064e-01	7.64024
3	9.999789e-01	4.743223e-08	7.266109e-08	9.39932
4	4.452492e-08	9.999914e-01	4.315461e-08	7.71225
...
106	9.999990e-01	4.436773e-08	4.735975e-08	6.70330
107	9.997505e-01	1.521141e-05	1.925700e-04	6.56598
108	9.999776e-01	2.305937e-06	1.045074e-06	9.33498
109	9.997518e-01	1.508667e-06	4.464774e-05	1.36007
110	9.905429e-01	8.528121e-03	6.604245e-04	1.17113

111 rows × 6 columns

1 # Concat dependant Ancestry features to dataframe

2 dep_preds = pd.concat([Test_dataset['Ancestry'], df2], axis = 1)

1 # Add new

2 dep_preds['ClassPredName'] = dep_preds['ClassPredInd'].apply(lambda x: Trainin

1 # Validate Probability predictions dataframe

2 dep_preds.head()

	Ancestry	African_Probability	European_Probability	EastAsian_Probability	Oceanian_Probability
0	Unknown	2.653985e-08	4.076232e-07	4.195787e-08	
1	Unknown	3.428250e-08	3.127784e-05	2.568485e-06	
2	Unknown	3.568520e-04	2.094020e-02	9.767064e-01	
3	Unknown	9.999789e-01	4.743223e-08	7.266109e-08	
4	Unknown	4.452492e-08	9.999914e-01	4.315461e-08	

1 # Slice prediction and set new feature vector column variable

2 prob_1 = dep_preds.loc[:, 'Ancestry':'NativeAmerican_Probability']

```

1 # Unpivot convert dataframe to long format
2 prob_2 = pd.melt(prob_1, id_vars = ['Ancestry'], var_name = 'Ancestry_Predicti

1 # Test for true probability
2 prob_2['Ancestry_Predictions'] = prob_2['Ancestry_Predictions'].apply(lambda x

1 # Validate dataframe
2 prob_2.head(5)

```

	Ancestry	Ancestry_Predictions	Probability
0	Unknown	African_	2.653985e-08
1	Unknown	African_	3.428250e-08
2	Unknown	African_	3.568520e-04
3	Unknown	African_	9.999789e-01
4	Unknown	African_	4.452492e-08

```

1 # Validate dataframe features
2 print('Describe Columns:=', prob_2.columns, '\n')
3 print('Data Index values:=', prob_2.index, '\n')
4 print('Describe data:=', prob_2.describe(), '\n')

```

Describe Columns:= Index(['Ancestry', 'Ancestry_Predictions', 'Probability'], dtype='ob

Data Index values:= RangeIndex(start=0, stop=555, step=1)

```

Describe data:=          Probability
count  5.550000e+02
mean    2.000000e-01
std      3.713757e-01
min      1.157843e-08
25%      2.527892e-05
50%      9.045926e-04
75%      9.444702e-02
max      9.999995e-01

```

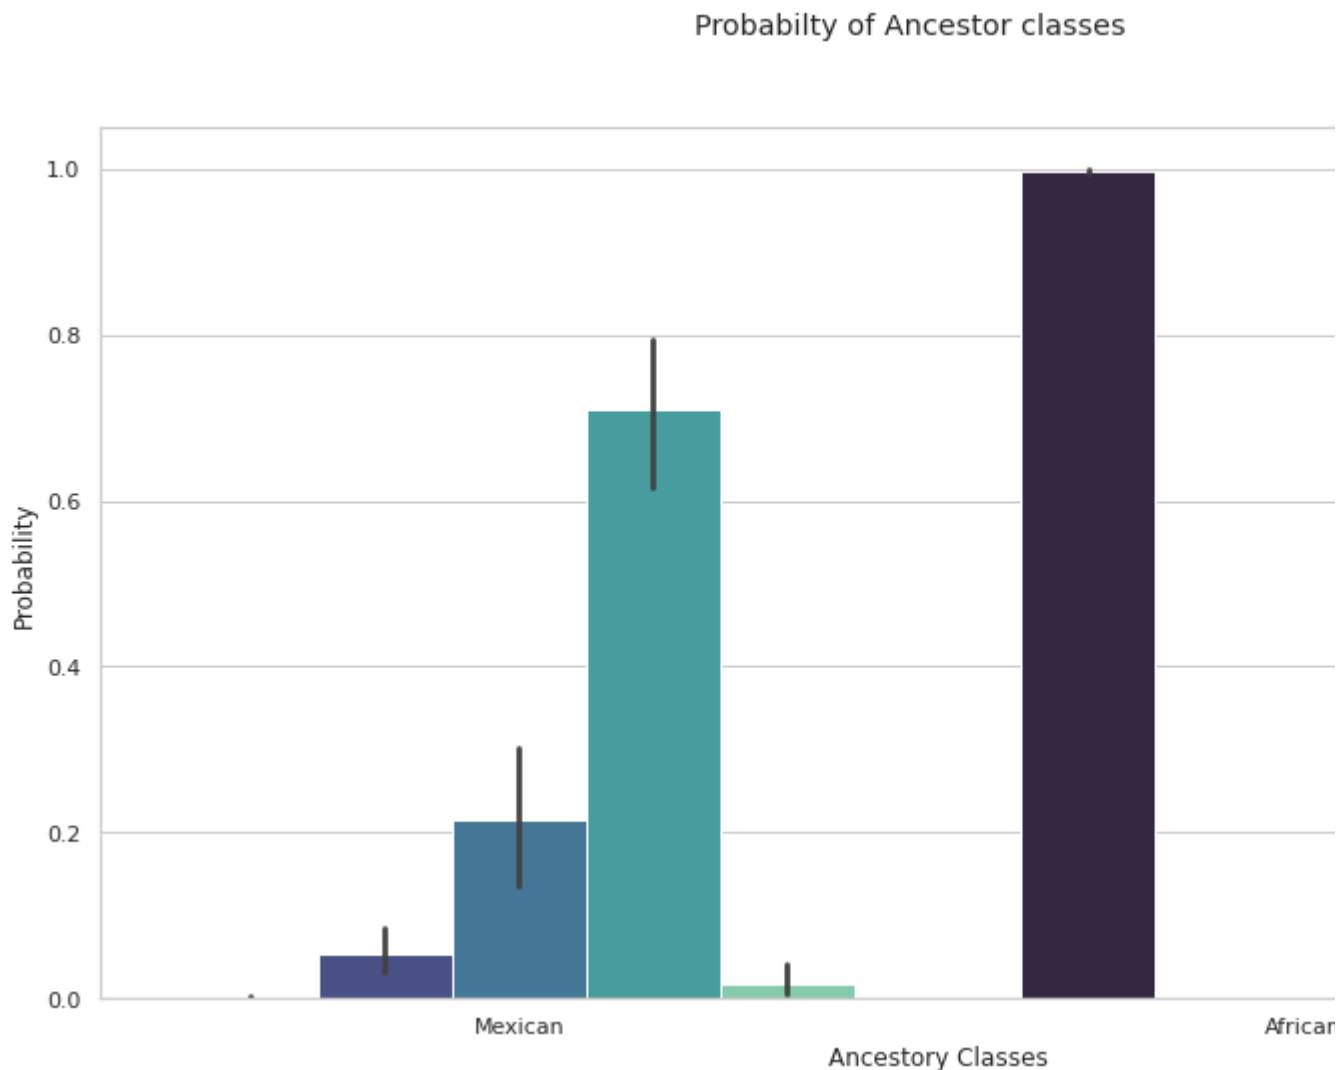


```

1 # Plot Probability prediction matrix
2 sns.set(rc = {'figure.figsize':(15,8)})
3 sns.set_theme(style="whitegrid")
4 fig, ax = plt.subplots()
5 sns.barplot(data = prob_2[prob_2['Ancestry'] != 'Unknown'],color = 'r', x =
6 plt.xlabel('Ancestry Classes')
7 plt.ylabel('Probability')
8 plt.suptitle('Probability of Ancestor classes')

```

```
9 plt.savefig("Assignment3_Deliverable4.png")
10 plt.show()
```



▼ Deliverable 7.5

How do the class label probabilities differ for the Mexican and African American samples when compared to the class label probabilities for the unknown samples? Are these class probabilities telling us something about recent history? Explain why these class probabilities are reasonable with respect to knowledge of recent history?

- In comparison to the class label probabilities for the unknown samples, those with unknown ancestry show a probability close to or equal to one while the other classes show a probability close to zero or less than one. African American samples showed similar results. The model assigned high probabilities to the African ancestry class for each of these